

PITCH DECK · CONTENT BOOK

Margdarshan.ai

Multilingual AI Grievance Redressal — built for citizens, designed for governments

BGI Hackathon 2026 · Problem BT1P1

Team: Merge Conflict (Team ID 2515)

How to use this document. Each page below is one slide. The top half is what the audience sees (title + bullets + visual direction). The yellow box at the bottom is the speaker script — read aloud or paraphrase. The deck targets a **10-minute pitch + 5-minute Q&A**; per-slide timings are noted. Slide order is final; bullets are paste-ready.

Deck spec

Aspect ratio: 16:9 · **Font:** Inter / Helvetica / Calibri · **Primary:** #1E3A8A · **Accent:** #F59E0B · **Body text:** #0F172A on white. Use the same palette as the product UI.

DECK MAP

18 slides · 10 minutes + 5 Q&A

#	Slide	Time	Purpose
01	Title	0:15	Brand + hook
02	The Problem	0:45	Pain in numbers
03	Why Today's Tools Fail	0:45	Gap analysis
04	Introducing Margdarshan.ai	0:30	One-liner + tagline
05	Live Demo Flow	1:00	Citizen journey in 10 steps
06	System Architecture	0:45	Microservices map
07	AI/ML Stack	0:45	LLM + MuRIL + S-BERT
08	70-Portal Routing	0:30	Why hierarchical lookup wins
09	Semantic Deduplication	0:30	S-BERT cosine ≥ 0.85
10	Submission & Reliability	0:30	Adapter + 7-attempt retry
11	Adaptive Status Tracking	0:30	1h → 6h → 24h cadence
12	Government Dashboard	0:45	Review + audit
13	Multilingual UX & Voice	0:30	Hindi / English / Hinglish
14	Security, Audit & Trust	0:30	SHA-256 hash chain + JWT
15	Differentiators	0:30	vs CPGRAMS / 181 / MyGov
16	Impact, Scale & Roadmap	0:45	What unlocks at 1M citizens
17	Team Merge Conflict	0:15	Roles + credit
18	Thank You / Ask	0:15	Q&A handle

SLIDE 01

Margdarshan.ai

Guiding 1.4 billion citizens to the right grievance portal — in their own language

On-screen content

- Product name: **Margdarshan.ai**
- Tagline: “*Bolo, bot sunega — sahi jagah pahunchayega.*” (Speak — the bot listens — and routes you to the right place.)
- Problem statement: **BT1P1 — AI for Citizen Grievance Redressal**
- Team: **Merge Conflict** · Team ID **2515** · BGI Hackathon 2026

***VISUAL:** Full-bleed background — Indian flag tricolour gradient at 10% opacity; centred logo (chat bubble + compass needle). Bottom-left QR code linking to live demo at <http://localhost:5173>. Bottom-right team photo strip.*

SPEAKER NOTES & DELIVERY

Open warm and human. “**Imagine a labourer in Bhopal whose electricity has been out for three days. He doesn't speak English, doesn't know which of seventy government portals to file at, doesn't even own a smartphone comfortable typing Devanagari. Today, his complaint dies in a queue. Tonight, we change that.**” Pause. Then introduce the team in one breath: “*We're Team Merge Conflict — and this is Margdarshan.ai.*”

Timing: 15 seconds. **Energy:** calm, confident — do not rush.

SLIDE 02

The Problem

Grievance redressal in India is broken at every step

70+ government portals fragmented, overlapping mandates	47 days avg resolution time CPGRAMS public data, 2024	~38% complaints filed wrong portal central vs state mismatch
--	--	---

The five fault lines

- **Language barrier:** Most portals are English-first; 78% of India is more comfortable in a regional language.
- **Portal maze:** The same complaint (e.g., bijli) routes differently per state — CPGRAMS, MP CM Helpline 181, MPPKVCL, BSES, ...
- **Duplicates:** 30 households reporting one transformer fault create 30 tickets — wasting officer time.
- **No closure loop:** Citizens don't know if it was resolved unless they keep checking the portal.
- **No accountability:** No tamper-evident record of who saw the complaint and when.

VISUAL: Left half: stacked-bar chart — share of complaints filed at the wrong portal (38%) vs right portal (62%). Right half: photo of a queue at a government office, low saturation.

SPEAKER NOTES & DELIVERY

Anchor with the three KPI cards. **“Forty-seven days. That’s the average. For an electricity outage. Thirty-eight percent of complaints don’t even land at the right office — they bounce. And there are over seventy portals to choose from.”** Pause on the number 70. That’s the wedge.

Sources to cite if asked: CPGRAMS annual report 2024; DARPG dashboards; MyGov 2023 stats. Don’t put citations on the slide — keep them in your head.

SLIDE 03

Why Today's Tools Fail

Forms, helplines, and chatbots each miss something critical

	CPGRAMS web form	MP CM Helpline 181	MyGov chatbot	Margdarshan.ai
Native Hindi / Hinglish input	Partial	Voice only	No	Yes (text + voice)
Auto-routes to correct portal	No	Manual operator	No	Yes (70 portals)
De-duplicates similar complaints	No	No	No	Yes (S-BERT, cosine ≥ 0.85)
Adaptive status polling	No	Manual callback	No	Yes (1h → 6h → 24h)
Tamper-evident audit trail	No	No	No	Yes (SHA-256 hash chain)
Department dashboard with AI review	No	Limited	No	Yes (override + retrain)

VISUAL: Highlight the rightmost column ('Margdarshan.ai') with the accent yellow row stripe. Animate row-by-row reveal on click.

SPEAKER NOTES & DELIVERY

“CPGRAMS is a form. 181 is a helpline. MyGov is a chatbot. None of them put the entire pipeline together. We did.” Walk down the right column row-by-row — voice, routing, dedup, polling, audit, dashboard. Each one is a feature competitors don't have.

SLIDE 04

Introducing Margdarshan.ai

A conversational pipeline — speak the problem, we do the rest

Elevator pitch (read aloud verbatim)

“Margdarshan.ai is a multilingual AI grievance pipeline that lets any citizen describe a problem in Hindi, English, or Hinglish — by typing or speaking. We classify the complaint into one of 25 departments and 142 sub-categories, deduplicate against existing tickets using multilingual S-BERT, route to the correct portal out of more than seventy, and track resolution end-to-end with a tamper-evident audit chain. The same backend powers a department dashboard where reviewers can override AI decisions, and those overrides become training data for the next model cycle.”

Product surfaces

- **Citizen UI** — chat (typed + voice), location pin, bilingual notifications.
- **Department Dashboard** — stats, drill-down, override, duplicate clusters, audit timeline.
- **REST API** — gateway exposes /chat, /complaints, /admin/* — pluggable by other govt apps.

***VISUAL:** Side-by-side mockups: phone screen (citizen chat in Hinglish) | laptop screen (dashboard stats). Connect them with a centred arrow labelled ‘same backend, same data, two audiences’.*

SPEAKER NOTES & DELIVERY

Land the one-liner cleanly. After reading the pitch, pause two seconds. Then: **“One pipeline. Two audiences. Built in eight weeks.”** Move on — don't dwell.

SLIDE 05

Live Demo — The Citizen Journey

One complaint, ten stages, ninety seconds

Stages (numbered for live narration)

- **1. Chat:** “bijli 3 din se nahi aa rahi, Bhopal MP” — typed or spoken (Web Speech API, hi-IN).
- **2. NLU:** Claude Haiku 4.5 extracts intent (COMPLAINT_NEW), entities (bijli, 3 din), language (Hinglish).
- **3. Classify:** MuRIL → dept ELECTRICITY (0.94), sub-cat POWER_OUTAGE (0.88), priority HIGH, sentiment FRUSTRATED.
- **4. Location:** Leaflet map modal; LocationIQ reverse-geocode → Bhopal, pincode 462001.
- **5. Route:** 70-portal lookup — Regional > State > Central; matches MPPKVVCL (regional discom).
- **6. Dedup:** S-BERT cosine vs last 30 days, district-scoped; threshold 0.85; new complaint, not dup.
- **7. Field collection:** portal asks for service-number; Claude prefills location/description; mobile regex validates.
- **8. Submit:** MPPKVVCL adapter POSTs; ticket id **ELE/MPB/0042** returned in 0.8s.
- **9. Track:** 1h poll cadence begins; status normaliser handles portal vocabulary.
- **10. Close loop:** WhatsApp (Hindi → English) on every status change; on RESOLVED, ‘1=satisfied, 2=reopen’.

***VISUAL:** Animated 10-step horizontal timeline across the bottom; live screen-mirror of the citizen UI on the top half. If demo Wi-Fi fails, fall back to a 30-second pre-recorded mp4 — keep it on the speaker's laptop.*

SPEAKER NOTES & DELIVERY

Live demo strategy: open with the spoken Hinglish complaint (microphone on). Let the audience hear your voice go into the bot. After step 5, switch to the dashboard tab to show the same complaint appearing live.

Always have the recording as backup. Wi-Fi at hackathons is unreliable — don't risk it.

Talking point: emphasise that steps 2-10 happen automatically — the citizen only does step 1 and pins location.

SLIDE 06

System Architecture

Ten microservices on Docker Compose — independently deployable

Services (port → responsibility)

Service	Port	Responsibility
web-frontend	5173	React 18 + Vite + Tailwind — citizen chat + admin dashboard
service-gateway	8000	Public REST + WebSocket · JWT · pipeline orchestrator · SHA-256 audit chain
service-chatbot	8001	Thin chat interface · Claude-driven field-collection mode
service-location	8002	LocationIQ + OSM Nominatim · India Post pincode
service-nlu	8003	State machine · intent · entities · language detect · Claude decision JSON
service-classifier	8004	MuRIL fine-tuned + 4 sklearn heads (rule-based fallback toggle)
service-routing	8005	70-portal hierarchical lookup (Regional & State & Central)
service-submission	8006	Adapters: CPGRAMS / MPCM181 / MPPKVVCL · 7-attempt backoff
service-tracker	8007	Adaptive polling · status normaliser · mock Twilio WhatsApp
postgres / redis	—	JSONB UCO + FLOAT8[] vectors · Redis pub/sub on pipeline:all

VISUAL: Render the architecture diagram from the README (gateway in the centre, six worker services radiating out, tracker below, Postgres + Redis at the bottom). Use the boxes in BRAND_PRIMARY with white text.

SPEAKER NOTES & DELIVERY

Don't read the table. Point at three things: **(1)** Gateway is the only public face — everything else is internal. **(2)** Redis pub/sub is the event bus — every stage emits a PipelineEvent, the gateway orchestrates. **(3)** Each service has its own Dockerfile and requirements.txt — we can independently scale the classifier without touching the chatbot.

SLIDE 07

AI / ML Stack

Three models, each chosen for a different job

Claude Haiku 4.5 Conversational LLM ~700 ms · temp 0.2	MuRIL fine-tuned Classifier (dept) F1 0.91 · 25 classes	S-BERT multilingual Dedup embedding 384-dim · cosine ≥ 0.85
--	---	---

Why this split — and not one big LLM for everything

- **Latency budget.** Claude Haiku for dialogue (we need it to feel like chat). MuRIL for classification (deterministic, ~30 ms, no API cost).
- **Reliability.** Rule-based classifier as a hot fallback (USE_RULE_BASED_CLASSIFIER=true) — no model server, no GPU dependency.
- **Multilingual coverage.** MuRIL was pre-trained on 17 Indian languages — better Hindi/Hinglish than English-first BERT/DistilBERT.
- **Confidence-aware UX.** If classifier confidence < threshold, the chatbot asks a clarifying question — never silently mis-routes.
- **Human-in-loop.** Dashboard reviewers can override; overrides land in complaint_reviews → next training cycle.

VISUAL: Three-card row repeated visually as 'one model per job' — each card a different colour: blue (Haiku), green (MuRIL), amber (S-BERT). Below: a small horizontal latency bar — Haiku ~700ms, MuRIL ~30ms, S-BERT ~50ms — total pipeline under 1s.

SPEAKER NOTES & DELIVERY

Anticipated judge question: 'Why not just use one LLM call for everything?' Answer: **cost, latency, and determinism.** Classification needs to be reproducible — a tuned MuRIL model is auditable. Dedup needs vector math — embeddings are the right primitive. Use the LLM where conversation matters, not where math suffices.

SLIDE 08

70-Portal Routing

Hierarchical specificity — never a central catch-all when a regional desk exists

The lookup hierarchy

- **1. Regional** — e.g., MPPKVVCL for Bhopal electricity (matches dept + district).
- **2. State-specific** — e.g., MP CM Helpline 181 (matches dept + state).
- **3. State catch-all** — e.g., state government CMO portal.
- **4. Central-specific** — e.g., Railways grievance.
- **5. Central catch-all** — CPGRAMS as last resort.

Why it matters

A power cut in Bhopal sent to CPGRAMS will sit in a national queue for weeks before being forwarded to the discom. Sent directly to MPPKVVCL, the discom owns it immediately. The routing service implements a **specificity-sorted match** across portals.csv — 70 entries today, designed to scale to 200+ without code changes.

VISUAL: Show a 5-tier pyramid (regional at top, CPGRAMS at base). Animate a 'Bhopal electricity' chip falling and stopping at tier 1. Then a second chip 'pan-India IT grievance' falling to tier 4.

SPEAKER NOTES & DELIVERY

“We didn't just digitise CPGRAMS. We made CPGRAMS the fallback, not the default.” That line lands every time. Then walk through the Bhopal example — concrete, visual, ends with a portal name the judges may not have heard of.

SLIDE 09

Semantic Deduplication

30 households, 1 transformer fault, 1 ticket

How it works

- Every new complaint is embedded with **paraphrase-multilingual-MiniLM-L12-v2** (384-dim, 117 MB, runs locally).
- Cosine similarity is computed against complaints from the **same district in the last 30 days** (partial index on Postgres FLOAT8[]).
- ≥ 0.85 → flagged as duplicate; linked to the parent ticket; dashboard surfaces the cluster.
- Multilingual: “*bijli nahi aa rahi*” and “*no electricity since yesterday*” embed to vectors with cosine ~ 0.91 .

Why this is non-trivial

- Keyword dedup would miss cross-language duplicates entirely.
- Full-text search misses paraphrases (‘power cut’ vs ‘outage’).
- S-BERT handles both — and the multilingual variant means we don't run two models for two languages.

VISUAL: Cluster diagram: 30 small dots with cosine-similarity values labelled, all connected to one central 'parent' dot. Threshold line drawn at 0.85. Annotate one cross-language pair in red.

SPEAKER NOTES & DELIVERY

“This is the feature officers love most when we demo it.” Reviewers spend the bulk of their time deduping. Anchor on the transformer-fault example. Mention: it's not just user-level dedup — it's also citizen-vs-citizen clustering, which CPGRAMS can't do.

SLIDE 10

Submission & Reliability

Plug-in adapter pattern + exponential-backoff retries

Adapter contract (4 methods)

submit(uco) → ticket_id · check_status(ticket_id) → status · list_fields() → schema · validate(payload) → errors[]

Adapters shipped

- **CPGRAMSAdapter** — central catch-all (mock POST; real schema modelled).
- **MPCM181Adapter** — MP CM Helpline 181 (state).
- **MPPKVVCLAdapter** — Madhya Pradesh discom (regional).
- **GenericAdapter** — fallback when no portal-specific adapter exists.

Retry schedule (per complaint)

7 attempts, exponential backoff: 1 min → 2 min → 5 min → 15 min → 1 h → 4 h → 24 h. After the 7th failure, the complaint enters the dashboard's 'manual intervention' queue.

***VISUAL:** Timeline graphic — the seven backoff steps as widening boxes. Below each, a small icon (network blip, portal maintenance, etc.) showing what each interval absorbs.*

SPEAKER NOTES & DELIVERY

Why this matters: government portals go down. We've watched CPGRAMS return 502s for two-hour windows. Our 7-attempt curve absorbs that and any single-day outage. After that, a human picks it up — never silently dropped.

SLIDE 11

Adaptive Status Tracking

Polls hard early — relaxes once stable

Polling cadence

- **First 24 hours:** every 1 hour (status changes most often early).
- **Day 2–7:** every 6 hours.
- **Day 8+:** every 24 hours.
- **Stops** on terminal state (RESOLVED / REJECTED).

Status normaliser

Portals use different vocabularies — 'Under Process', 'Pending with department', 'In progress', 'Action taken' all mean the same thing. The normaliser maps every portal-specific string to one of: **SUBMITTED · UNDER_REVIEW · IN_PROGRESS · RESOLVED · REJECTED · ESCALATED**.

Citizen notifications

- On every status change, WhatsApp (mock Twilio) message — **Hindi first, English second**.
- On RESOLVED, follow-up prompt: “1 = satisfied · 2 = reopen”. Reply '2' re-files with the original UCO.
- Feedback closes the loop; satisfaction rate is a dashboard metric.

***VISUAL:** Stacked line chart — polling frequency drops from hourly to daily over a 30-day window. Below it, a phone-screen mockup of the bilingual WhatsApp notification.*

SPEAKER NOTES & DELIVERY

Killer line: ‘The citizen never has to log in again to check status. The system tells them — in Hindi — when something changes.’

SLIDE 12

Government Dashboard

The same backend, a different audience

Eight surfaces

- **Stats overview** — hero cards (total / duplicates / by status / by dept / by district / 10 most recent).
- **Portal registry** — all 70 portals with helpline, authority, complaint count routed.
- **Complaints list** — filter by status, dept, district; priority and sentiment badges.
- **Complaint drill-down** — filer info (PII-masked), top-3 classification alternatives, routing explanation, full audit timeline.
- **Review form** — thumbs up / down, override department, sub-cat, priority, sentiment → complaint_reviews table.
- **Duplicate filers** — repeat filers + cross-user clusters (the 30-household transformer case).
- **Routing explanation** — exact rule that fired (tier, district match, fallback path).
- **API trace** — animated 5-step demo view showing live JSON between services.

***VISUAL:** Stitched screenshot collage — stats hero (top-left), complaint drill-down (top-right), duplicate cluster view (bottom-left), review form (bottom-right). Use the real production-tier dashboard screens.*

SPEAKER NOTES & DELIVERY

Switch to the dashboard tab if you haven't already. Click into a complaint with an override pending — show the thumbs-down flow. **“Every override here is training data tomorrow. The model gets better the more it's used.”**

SLIDE 13

Multilingual UX & Voice

Speak. Type. Mix. We hear you.

Languages — UI strings, bot replies, notifications

- ■■■■■■ (Devanagari) — full UI + bot replies + WhatsApp notifications.
- **English** — same coverage.
- **Hinglish** (Roman-script Hindi) — default for new users; the most common comfort-language in tier-2/3 India.
- Switchable any time from the top-right language selector; preference persists per user.

Voice input

- Browser Web Speech API: locales hi-IN and en-IN.
- Tap the mic icon → speak the complaint → bot transcribes → pipeline runs.
- Tested on Chrome Android (the dominant browser-OS combo in India).

Auto-detect

Even before language preference is set, the NLU layer detects script + tokens — Hindi vowels vs Roman characters vs mixed — so the bot replies in the same language the citizen used, not in the language they once configured.

VISUAL: Three phone-screen mockups side-by-side — same complaint, three languages. Show the bot reply matching the input language.

SPEAKER NOTES & DELIVERY

Pick up the mic on stage. Switch the laptop to phone-emulator view. Speak the complaint **in Hindi** live: “Bhopal mein bijli nahi aa rahi, teen din se.” Watch the room when the bot replies in Hindi.

SLIDE 14

Security, Audit & Trust

Tamper-evident by construction — not by policy

Audit chain

Every state change for a complaint is recorded in `complaint_events` with a **SHA-256 hash of the previous event**. Any retroactive tampering breaks the chain. Dashboard shows the chain visually — a single broken link is visible to any reviewer.

Authentication

- bcrypt password hashing (cost 12).
- JWT (HS256), 24-hour expiry, stored in `localStorage`; refresh on activity.
- Role-based access: citizen, reviewer, admin.

PII handling

- Mobile, Aadhaar, and email masked in the dashboard list view (last-4 only).
- Aadhaar field uses **Verhoeff checksum validation** — invalid Aadhaars rejected at input, never stored.
- Full PII visible only on drill-down behind reviewer role.

Field validation (11 regex validators)

pincode (6 digits) · mobile (10 digits, prefix 6-9) · email · Aadhaar (Verhoeff) · vehicle reg · service number · consumer number · IFSC · GSTIN · date · time. Validation is portal-aware — each portal declares its required fields.

SPEAKER NOTES & DELIVERY

Anticipated judge question: 'What about data privacy?' Answer: PII is masked by default; full unmask is role-gated; nothing leaves our infra except the portal POST (which is the citizen's intent).

SLIDE 15

What Sets Us Apart

Six things we have that the alternatives don't

- **True multilingual + voice input** from the first message — not bolted on as a translation layer.
- **70-portal hierarchical routing** — regional desks beat central catch-all every time.
- **Multilingual S-BERT dedup** across citizens, not just per-user.
- **Adapter pattern** — adding a new portal is a single Python class implementing four methods; no core changes.
- **Adaptive polling + bilingual notifications** — the citizen never has to come back to check.
- **SHA-256 audit chain + reviewer override loop** — every override is training data for the next model.

One-line moat

“We're not a chatbot in front of a form. We're a full pipeline — and the form is just the last 10% of it.”

VISUAL: Six icon-tile grid — one tile per differentiator, with the icon in *BRAND_ACCENT*. Bottom band: the italic one-liner in 22pt, centred.

SPEAKER NOTES & DELIVERY

If you only have one minute and have to pick one differentiator: **say the dedup story**. It's the most concrete and the easiest to visualise. Save the audit chain for the security-question follow-up.

SLIDE 16

Impact, Scale & Roadmap

From hackathon prototype to national infrastructure

Today (hackathon scope)

- 70 portals · 25 departments · 142 sub-categories · 3 languages.
- End-to-end pipeline live on Docker Compose — single host.
- Mocked outbound POST to portals (no MoUs yet); everything else is real.

Next 3 months

- Pilot with one state (target: Madhya Pradesh) — real CPGRAMS + 181 + MPPKVVCL POSTs.
- Voice IVR fallback — phone number for citizens without smartphones.
- Active learning loop — auto-retrain MuRIL weekly on reviewer-overridden labels.

Next 12 months

- Expand to 200 portals across 8 states.
- Kubernetes deploy (HPA on classifier & gateway); RPS target 500.
- Open the API to other citizen-services apps (file an FIR, book an ambulance — same pipeline shape).

Why this scales

Microservices independently scalable; classifier is CPU-only (no GPU dependency in the rule-based mode); Redis pub/sub is horizontally shardable; Postgres FLOAT8[] indices are pgvector-compatible if we outgrow them. Cost per complaint today: **under ₹0.30** (Claude Haiku tokens dominate).

SPEAKER NOTES & DELIVERY

Judge-mind read: they'll ask about scale and unit economics. Have the ₹0.30/complaint number ready. Have the 500 RPS number ready. Have the 'one new portal = one Python class' number ready.

SLIDE 17

Team Merge Conflict

Team ID 2515 · BGI Hackathon 2026 · BT1P1

Roles (fill in names)

- **Backend / Pipeline:** [Name] — services, gateway orchestrator, audit chain, adapter pattern.
- **AI / ML:** [Name] — MuRIL fine-tune, S-BERT dedup, classifier rule-set, Claude prompt design.
- **Frontend:** [Name] — React + Vite SPA, Leaflet, i18n, dashboard, citizen chat UX.
- **DevOps / Data:** [Name] — Docker Compose, Postgres schema, training dataset, portals.csv curation.

Build stats

10 services shipped 9 backend + 1 web	70 portals onboarded regional → central	3 languages Hindi · English · Hinglish
VISUAL: Photo strip of the four team members (square headshots) above the role bullets. GitHub handles + LinkedIn QR codes at the bottom-right corner.		
SPEAKER NOTES & DELIVERY Quick. One sentence per member: what they built, not their CV. “[Name] built the entire submission and retry layer in 72 hours.” That kind of line — concrete, accomplishment-anchored.		

SLIDE 18

Thank You

Margdarshan.ai — bolo, bot sunega

The ask

- Pilot opportunity with one state government department.
- Introduction to DARPG / state CMOs for CPGRAMS-compatible MoU.
- Mentorship on government procurement & data-governance compliance (DPDP Act).

Try it live

- Citizen UI: **<http://localhost:5173>**
- Gateway API docs: **<http://localhost:8000/docs>**
- Admin login: **[admin@margdarshan.ai](#)** / **[Admin@1234](#)**
- GitHub: [team repo URL] · Technical doc: *docs/Margdarshan-ai-Technical-Doc.pdf* (35 pages)

VISUAL: Full-bleed background — same as the title slide. Big centred: 'Q&A'. Bottom row: four small icons — 'pilot', 'mentorship', 'introductions', 'feedback' — each clickable in the slide deck.

SPEAKER NOTES & DELIVERY

Closing line — say it slowly: "We started this with one belief — that a labourer in Bhopal should not need English, a laptop, or a lawyer to reach his government. Margdarshan.ai is our first attempt. Thank you." Pause. Take questions.

Q&A prep: have ready answers for — (1) cost/complaint, (2) privacy & DPDP, (3) what's real vs mocked, (4) how a new portal gets added, (5) failure modes & the manual queue, (6) accuracy numbers (dept F1 0.91, sub-cat 0.85).

APPENDIX

Design system + delivery checklist

Colour palette (hex)

Role	Hex	Usage
Primary (deep blue)	#1E3A8A	Headings, primary buttons, brand bars
Accent (amber)	#F59E0B	Highlight rows, CTA buttons, slide rules
Dark text	#0F172A	Body copy on white
Mid grey	#475569	Subtitles, captions
Card background	#EFF6FF	KPI cards, info boxes
Note background	#FEF3C7	Speaker-note callout boxes
Success green	#15803D	Status RESOLVED, success indicators
Error red	#B91C1C	Negative metrics, warnings

Typography

- **Primary font:** Inter (web) / Helvetica (fallback) / Calibri (PowerPoint default — fine).
- **Slide title:** 36-44pt bold.
- **Section heading:** 20-24pt semibold.
- **Body bullets:** 18-22pt regular.
- **Captions / legends:** 12-14pt regular italic.
- **Monospace** (code/IDs): JetBrains Mono / Menlo / Courier.

Slide template (16:9, master layout)

- Top bar (5% height) — brand primary fill; left = 'MARGDARSHAN.AI' in white bold; right = slide number.
- Content area (90%) — white background; 60px left/right margin; 40px top/bottom.
- Footer (5%) — light grey, centred page indicator; bottom-right discreet 'Team Merge Conflict · 2515'.
- Accent rule (2px, amber) under each slide title.

Iconography

- Use **Lucide** or **Heroicons** outline set — consistent 1.5px stroke.
- Icon size in slides: 32-40px; tint in BRAND_PRIMARY by default, BRAND_ACCENT for emphasis.
- Avoid emoji on judging slides; emoji fine for backup notes.

Animations (Keynote / PowerPoint)

- Slide transition: **Cross-fade, 0.3s** — anything fancier looks cheap on a projector.
- Bullet reveal: **Fade-in on click** — never auto-advance bullets.
- Demo slide (#5): pre-record a 30s mp4 as the background; overlay live cursor when network works.

APPENDIX

Delivery checklist — the 10 hours before the pitch

T-minus 10 hours

- Build the final docker stack: docker compose up -d --build.
- Run bash scripts/smoke_test.sh — must be green end-to-end.
- Verify admin login works; verify at least 5 distinct demo complaints exist with varied statuses.
- Pre-record a 30-second mp4 of the live citizen flow as Wi-Fi backup. Store on speaker's laptop, USB stick, and Google Drive.

T-minus 4 hours

- Rehearse the 10-minute pitch end-to-end **three times** on the actual deck.
- Time each slide. Cut anything that pushes past 10:00.
- Charge laptops to 100%; pack two HDMI/USB-C dongles.
- Print one A4 of the deck map + speaker notes per team member (paper backup if laptop dies).

T-minus 1 hour

- Disable laptop notifications, screensaver, low-power mode, OS update banners.
- Open the deck, the dashboard tab, the citizen chat tab — in that order; pre-pin.
- Plug into venue projector; verify aspect ratio is 16:9; verify colours don't wash out.
- Decide who owns the keyboard, who owns the mic, who owns the laptop swap if needed.

On stage

- Project confidence: speak to the back row, not the laptop.
- If a demo step fails — say **'we'll skip to the recording'**; never debug in front of judges.
- Reserve last 30 seconds for the closing line on Slide 18 — do not skip it.
- Q&A: if you don't know an answer, say **'great question — we don't have that number with us; we can follow up by email'** — never bluff.

Likely judge questions — drilled answers

Q	A
What's real vs mocked?	Everything is real except the outbound POST to each government portal (we don't have MoU
How do you handle privacy / DPDP Act compliance?	PII masked by default; full unmask is role-gated for reviewers; no PII leaves our infra except
Cost per complaint?	Under ₹0.30 today, dominated by Claude Haiku tokens (~600 in, ~150 out per turn). Classifie
How accurate is the classifier?	MuRIL fine-tuned: dept F1 0.91 (25 classes), sub-cat ~0.85 (142 classes). Low-confidence ca
How does a new portal get added?	One Python class with four methods (submit, check_status, list_fields, validate) + one row in p
What happens if all 7 retries fail?	The complaint moves to a 'manual intervention' queue visible in the dashboard. A reviewer ca
How does this differ from MyGov chatbot or CPGRAMS?	MyGov is conversation only — no routing, no dedup, no submission, no tracking. CPGRAMS
Can the model retrain itself?	Yes. Every reviewer override lands in complaint_reviews. The training script consumes that ta

APPENDIX

Hard numbers to memorise

Metric	Value	Source
Government portals routed	70	portals.csv registry
Department classes	25	classifier head 1
Sub-category classes	142	classifier head 2
Departments F1 (MuRIL)	0.91	training eval set
Sub-category F1 (MuRIL)	~0.85	training eval set
Languages supported	3	Hindi · English · Hinglish
Dedup threshold (cosine)	0.85	S-BERT MiniLM-L12
Embedding model size	117 MB / 384-dim	paraphrase-multilingual-MiniLM-L12-v2
LLM latency	~700 ms	Claude Haiku 4.5, temp 0.2
Submission retry attempts	7	1m → 2m → 5m → 15m → 1h → 4h → 24h
Polling cadence (day 1)	1 hour	tracker schedule
Polling cadence (week 1)	6 hours	tracker schedule
Polling cadence (after)	24 hours	tracker schedule
Cost per complaint	< ₹0.30	Haiku tokens dominate
Microservices in stack	10	9 backend + frontend
Audit hash function	SHA-256	complaint_events chain
Field validators	11	pincode · mobile · Aadhaar · ...
Adapters shipped	4	CPGRAMS · MPCM181 · MPPKVVCL · Generic

One-line elevator pitch (memorise verbatim)

“Margdarshan.ai turns ‘bijli nahi aa rahi’ — typed or spoken in any Indian language — into a tracked, deduplicated, audit-trailed complaint at exactly the right government portal out of seventy. End to end. Under a second.”